



kaggle

Kaggle Winner Presentation (5th)

NeurIPS 2024 - Predict New Medicines with BELKA

<https://www.kaggle.com/competitions/leash-BELKA>

HENG CHER KENG

<https://www.kaggle.com/hengck23>

Agenda

1. Background
2. Summary
3. Feature selection & engineering
4. Training methods
5. Important findings
6. Simple model
7. Code review

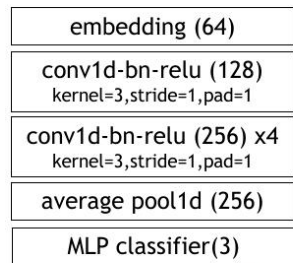
Background

- ◆ Contract computer vision and deep learning algorithm engineer.
 - identify fracture in x-ray images
 - implement visual slam for robotic navigation
 - finetune LLM models
- ◆ Familiar with deep learning and build deep models in my work.
- ◆ Experiences previous Kaggle competitions. Sequence modeling in molecule and medicine applications:
 - Bristol-Myers Squibb – Molecular Translation
 - OpenVaccine: COVID-19 mRNA Vaccine Degradation Prediction

Summary

cnn1d net

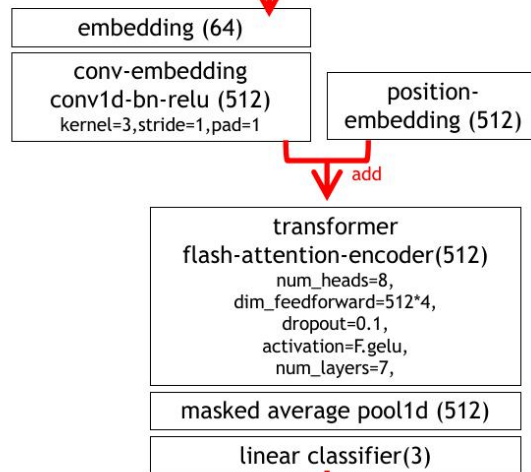
SMILES token id
(tokenised by character)



BRD4', HSA, sEH
(BCE loss)

transformer net

SMILES token id, token mask
(tokenised by character)



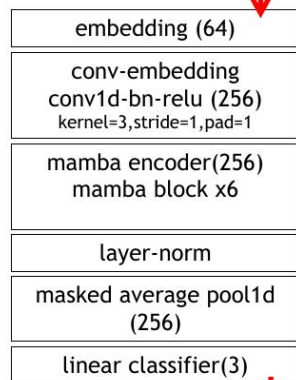
BRD4', HSA, sEH
(BCE loss)

◆ To mitigate the effects of OOD (out-of-domain distribution), we use ensemble of different sequence models:

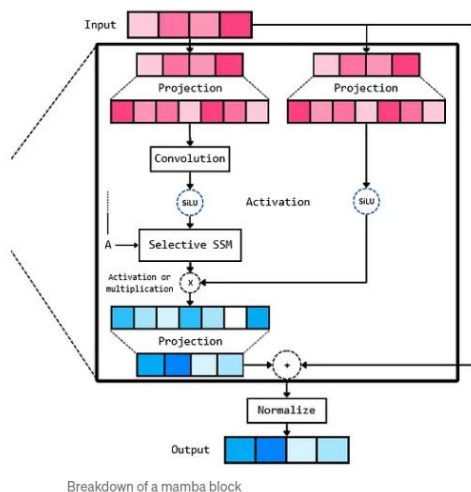
- 3-fold conv1d net
- 2-fold transformer net
- single-fold mamba (SSM, selective state machine)

◆ To handle 98 millions of training molecules efficiently, we use customized cuda kernel from open source flash attention[1] and mamba SSM[2] libraries

SMILES token id, token mask
(tokenised by character)



BRD4', HSA, sEH
(BCE loss)



◆ For training we have two Nvidia Ada A6000 / 48 GB Ampere GPUs. Time for training one fold with single GPU are:

- 7 hr for cnn1d,
- 28 hr for transformer,
- 36 hr for mamba.

			# keep score for the subset, set zero for others										
			all		keep-share #		keep-nonshare #		local CV (share)				
submission		fold	private	public	private	public	private	public	fold0	fold1	fold2	fold3	fold4
[1]+[2]+[3]	final-3fold-tx2a-mamba-fix.submit.csv		0.28557	0.46439	0.22761	0.34985	0.08320	0.10362	0.65976	0.66044		0.67601	
[1]	final-3fold-cnn1d.submit.csv	fold0,1,3	0.26615	0.41103	0.22693	0.34946	0.05424	0.08504					
[2]	final-2fold-transformer.submit.csv	fold2,4	0.31236	0.42288	0.22613	0.34753	0.10124	0.09881			0.64041		0.64589
[3]	final-1fold-mamba.submit.csv	fold0	0.23905	0.43221	0.22404	0.34484	0.03002	0.11083	0.65559				
late submission													
[2a]	final-2fold-cnn1d.submit.csv	fold2,4	0.26403	0.40879	0.22581	0.34814	0.05324	0.08411			0.63988		0.64613
[2b]	final-2fold-mamba.submit.csv	fold2,4	0.26361	0.41035	0.22712	0.34668	0.05150	0.08713			0.63848		0.64400

Features Selection/ Engineering

◆ Tokenization :

We tried several methods like character based, sentence piece, byte-pair-encoding (BPE), atom/smiles notation aware to break the SMILES strings. Surprisingly, the **simplest character based tokenization performs the best** across different net architecture

Add a conv1d layer of kernel size=3, stride=1 to learned combinations of consecutive tokens (bi-grams, tri-grams) before passing them into cnn1d, transformer or mamba encoder.

```
#https://www.ascii-code.com/
MOLECULE_DICT = {
    'l': 1, 'y': 2, '@': 3, '3': 4, 'H': 5, 'S': 6, 'F': 7, 'C': 8, 'r': 9, 's': 10, '/': 11, 'c': 12, 'o': 13,
    '+': 14, 'I': 15, '5': 16, '(': 17, '2': 18, ')': 19, '9': 20, '1': 21, '#': 22, '6': 23, '8': 24, '4': 25,
    '=': 26, '1': 27, '0': 28, '[': 29, 'D': 30, 'B': 31, ']' : 32, 'N': 33, '7': 34, 'n': 35, '-': 36
}
MAX_MOLECULE_ID = np.max(list(MOLECULE_DICT.values()))
VOCAB_SIZE=MAX_MOLECULE_ID+3
UNK=255 #disallowed, will cause error
BOS=MAX_MOLECULE_ID+1
EOS=MAX_MOLECULE_ID+2
PAD=0
MAX_LENGTH=160
```

◆ **Batch normalization:**

cnn1d model performance is sensitive to batch normalization. Different feature values for:

- in-distribution and out-distribution samples.
- +ve and -ve samples due to class imbalance (less than 1% +ve).

To alleviate the problem, we use high $\text{eps}=5\text{e-}3$ and low momentum=0.2

Training Methods

- ◆ binary cross entropy loss with ADAM optimizer.
 - step learning rate of 1e-3, 1e-4, 1e-5 for 6-12 epochs.
 - large batch size of 2000, 2500, 5000

Interesting, the best way to handle class imbalance is to do nothing (no up-sampling or under-sampling of the class). Maybe because of the large batch size we used.

Important and Interesting Findings

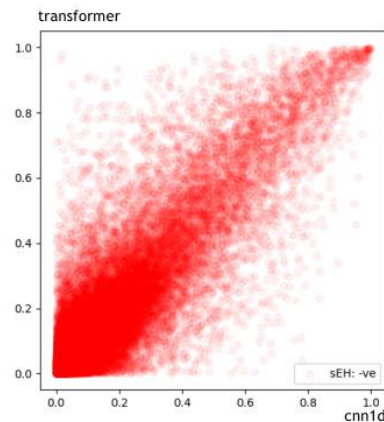
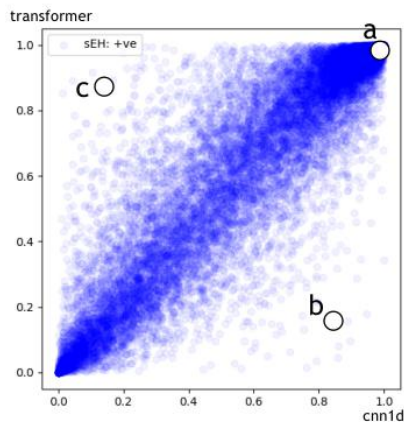
Important and Interesting Findings

◆ Transformer is the most robust

			# keep score for the subset, set zero for others										
			all		keep-share #		keep-nonshare #		local CV (share)				
submission		fold	private	public	private	public	private	public	fold0	fold1	fold2	fold3	fold4
[1]+[2]+[3]	final-3fold-tx2a-mamba-fix.submit.csv		0.28557	0.46439	0.22761	0.34985	0.08320	0.10362	0.65976	0.66044		0.67601	
[1]	final-3fold-cnn1d.submit.csv	fold0,1,3	0.26615	0.41103	0.22693	0.34946	0.05424	0.08504					
[2]	final-2fold-transformer.submit.csv	fold2,4	0.31236	0.42288	0.22613	0.34753	0.10124	0.09881			0.64041		0.64589
[3]	final-1fold-mamba.submit.csv	fold0	0.23905	0.43221	0.22404	0.34484	0.03002	0.11083	0.65559				
late submission													
[2a]	final-2fold-cnn1d.submit.csv	fold2,4	0.26403	0.40879	0.22581	0.34814	0.05324	0.08411			0.63988		0.64613
[2b]	final-2fold-mamba.submit.csv	fold2,4	0.26361	0.41035	0.22712	0.34668	0.05150	0.08713			0.63848		0.64400

◆ correlation of cnn1d and transformer predictions

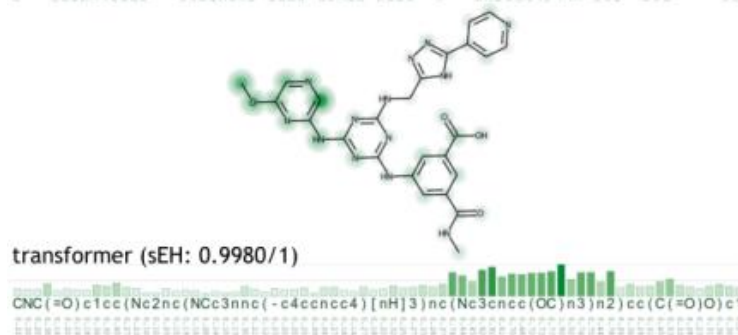
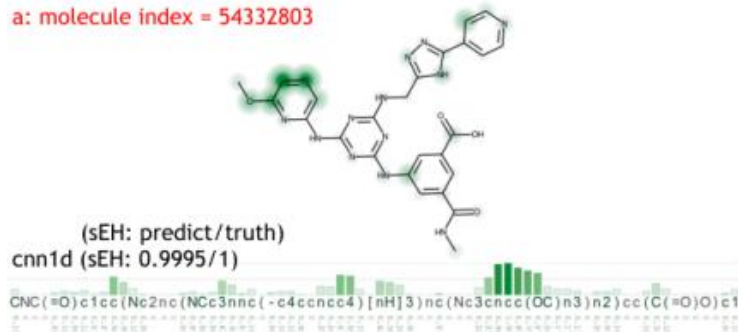
validation set: fold2 (4 million molecules)



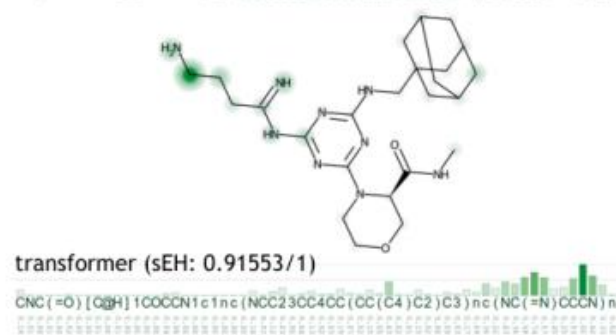
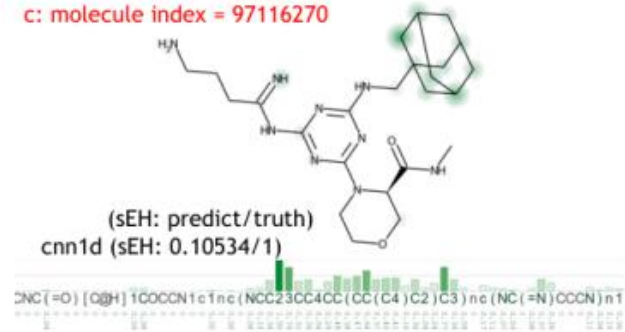
Important and Interesting Findings

- ◆ Generate the net prediction heatmap using GradCAM[3] and visualize it with XSMILES[4].
cnn1d has local activations, whereas transformer has more global ones.

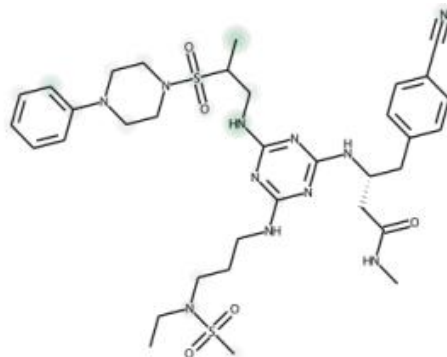
a: molecule index = 54332803



c: molecule index = 97116270



b: molecule index = 21821828



(sEH: predict/truth)
cnn1d (sEH: 0.9468/1)

CCN(CCCNc1nc(NCC(C)S(=O)(=O)N2CCN(c3ccccc3)CC2)nc(N[C@@H](C(=O)NC)Cc2ccc(C#N)cc2)n1)S(C)(=O)=O

transformer (sEH: 0.1779/1)
heatmap is all zeros

Simple Model

Simple Model

- ◆ Just transformer alone is good enough!

		# keep score for the subset, set zero for others										
		all		keep-share #		keep-nonshare #		local CV (share)				
submission	fold	private	public	private	public	private	public	fold0	fold1	fold2	fold3	fold4
[1]+[2]+[3] final-3fold-tx2a-mamba-fix.submit.csv		0.28557	0.46439	0.22761	0.34985	0.08320	0.10362	0.65976	0.66044		0.67601	
[1] final-3fold-cnn1d.submit.csv	fold0,1,3	0.26615	0.41103	0.22693	0.34946	0.05424	0.08504					
[2] final-2fold-transformer.submit.csv	fold2,4	0.31236	0.42288	0.22613	0.34753	0.10124	0.09881			0.64041		0.64589
[3] final-1fold-mamba.submit.csv	fold0	0.23905	0.43221	0.22404	0.34484	0.03002	0.11083	0.65559				
late submission												
[2a] final-2fold-cnn1d.submit.csv	fold2,4	0.26403	0.40879	0.22581	0.34814	0.05324	0.08411			0.63988		0.64613
[2b] final-2fold-mamba.submit.csv	fold2,4	0.26361	0.41035	0.22712	0.34668	0.05150	0.08713			0.63848		0.64400

Code Review

<https://github.com/hengck23/solution-leash-BELKA>

Kaggle Competition Solution (5th)

NeurIPS 2024 - Predict New Medicines with BELKA

<https://www.kaggle.com/competitions/leash-BELKA>

For discussion, please refer to:

<https://www.kaggle.com/competitions/leash-BELKA/discussion/456084>

1. Hardware

- GPU: 2x Nvidia Ada A6000 (Ampere), each with VRAM 48 GB
- CPU: Intel® Xeon(R) w7-3455 CPU @ 2.5GHz, 24 cores, 48 threads
- Memory: 256 GB RAM

2. OS

- ubuntu 22.04.4 LTS

3. Set Up Environment

- Install Python $\geq 3.10.9$
- Install requirements.txt in the python environment
- Set up the directory structure as shown below.

```
<solution_dir>
├── src
├── result
├── data
│   ├── processed
│   │   ├── all_buildingblock.csv
│   │   ├── kaggle
│   │   │   ├── leash-BELKA
│   │   │   │   ├── sample_submission.csv
│   │   │   │   ├── train.parquet
│   │   │   │   └── test.parquet
│   └── kaggle
│       ├── leash-BELKA
│       │   ├── sample_submission.csv
│       │   ├── train.parquet
│       │   └── test.parquet
├── LICENSE
└── README.md
```

- Download kaggle dataset "leash-BELKA" from:
<https://www.kaggle.com/competitions/leash-BELKA/data>

- Create processed data by run the python script:

```
python "/src/process-data-01/run_make_data.py"
```

There are 98 millions molecules in the train data. Hence processing the data can take very long time. Alternatively, you can download processed data from the share google drive at :
/leash-BELKA-solution/data/processed
https://drive.google.com/drive/folders/1bEBGtTJrQIYc_MQRYceBp0Kb9zGYue9H?usp=drive_link

- Modify the path setting by editing `"/src/third_party/_current_dir_.py"`

```
# please use full path
KAGGLE_DATA_DIR = '<solution_dir>/data/kaggle'
PROCESSED_DATA_DIR = '<solution_dir>/data/processed'
RESULT_DIR = '<solution_dir>/result'
```

4. Training the model

Warning !!! training output will be overwritten to the "/result" folder

Please run the following python scripts to learn the model files

```
python "/src/cnn1d-nonshare-05-mean-layer5-bn/run_train.py"
output model:
- /result/cnn1d-mean-pool-ly5-bn-01/fold-0/checkpoint/00400000.pth
- /result/cnn1d-mean-pool-ly5-bn-01/fold-1/checkpoint/00550000.pth
- /result/cnn1d-mean-pool-ly5-bn-01/fold-3/checkpoint/00415000.pth
```

```
python "/src/transformer-fa-03/run_train.py"
output model:
- /result/transformer-fa-03/fold-2/checkpoint/00264000.pth
- /result/transformer-fa-03/fold-4/checkpoint/00264000.pth
```

```
python "/src/mamba-03/run_train.py"
output model:
- /result/mamba-03/checkpoint/00255000.pth
```

If you want to do local validation, you can run the scripts:

```
python "/src/cnn1d-nonshare-05-mean-layer5-bn/run_valid.py"
python "/src/transformer-fa-03/run_valid.py"
python "/src/mamba-03/run_valid.py"
```

5. Submission csv

Please run the following scripts:

```
python "/src/cnn1d-nonshare-05-mean-layer5-bn/run_submit.py"
python "/src/transformer-fa-03/run_submit.py"
python "/src/mamba-03/run_submit.py"
python "/src/run_ensemble.py"
output file:
- /result/final-3fold-tx2a-mamba-fix.submit.csv
```

		# keep score for the subset, set zero for others					
		all		keep-share only #		keep-nonshare only #	
		private	public	private	public	private	public
	final-3fold-tx2a-mamba-fix.submit.csv	0.28557	0.46439	0.22761	0.34985	0.08320	0.10362
fold0,1,3	final-3fold-cnn1d.submit.csv	0.26615	0.41103	0.22693	0.34946	0.05424	0.08504
fold2,4	final-2fold-transformer.submit.csv	0.31236	0.42288	0.22613	0.34753	0.10124	0.09881
fold0	final-1fold-mamba.submit.csv	0.23905	0.43221	0.22404	0.34484	0.03002	0.11083

6. Reference trained models and validation results

- Reference results can also be found in the share google drive at :
/leash-BELKA-solution/result
https://drive.google.com/drive/folders/1bEBGtTJrQIYc_MQRYceBp0Kb9zGYue9H?usp=drive_link
- It includes the weight files, train/validation logs.

Authors

- <https://www.kaggle.com/hengck23>

License

- This project is licensed under the MIT License - see the [LICENSE](#) file for details.

Acknowledgement

"We extend our thanks to HP for providing the Z8 Fury-G5 Data Science Workstation, which empowered our deep learning experiments. The high computational power and large GPU memory enabled us to design our models swiftly."

Question and Answer



kaggle